

Week 4 Notes (Resume video at: --: 45: 19)

- Other Lifecycles:
 - o Recall: one reason to study OSes is that their abstraction/ideas/designs/techniques are relevant to many software projects, i.e., Android Lifecycle
 - Each app on Android goes through the same cycle
- A process can create processes
- If process A creates process B, "A is the **parent** of B", "B is the **child** of A"
 - o A process can have at most one parent
 - o A process can have many children
 - o Event handlers are responsible for the creation of processes
 - o At least one process needs to be running
- Each process has a **PID** (Process ID, integer)
 - o Picked by OS and is increasing
 - o i.e., if process created is PID 456, next process PID will be 457
 - o If PID is 2042, 2042 processes have been created since system boot
 - o C POSIX function "getpid()" returns PID
- Bottom line: Processes form a genealogy tree.
-

• On Mac OS X: ps axlw

```

UID  PID  PPID CPU PRI NI   VSZ  RSS WCHAN  STAT TT      TIME COMMAND
[...]
501  2660    1  0 31  0 2458784  536 -  Ss  ??  0:00.19 gpg-agent --daemon
501  2667    1  0 31  0 2467676  676 -  S  ??  0:00.08 /opt/X11/libexec/launchd_startx /opt/X11/bin/startx --
501  2668  2667  0 31  0 2439512 1064 -  S  ??  0:00.01 /bin/sh /opt/X11/bin/startx -- /opt/X11/bin/Xquartz
501  2733  2668  0 31  0 2452676  836 -  S  ??  0:00.08 /opt/X11/bin/xinit /opt/X11/lib/X11/xinit/xinitrc -- /c
501  2734  2733  0 31  0 2479128 2704 -  S  ??  0:00.01 /opt/X11/bin/Xquartz :0 -nolisten tcp -iglx -auth /User
501  2736  2734  0 63  0 2654600 46768 -  S  ??  0:06.31 /Applications/Utilities/XQuartz.app/Contents/MacOS/X11.
501  2743    1  0 31  0 2450592  532 -  Ss  ??  0:00.19 gpg-agent --daemon
501  2836  2733  0 31  0 2550224 7108 -  S  ??  0:00.07 /opt/X11/bin/quartz-wm
[...]

```

• On Linux: ps faux or ps --forest -eaf

```

UID  PID  PPID C STIME TTY      TIME  CMD
[...]
daemon 1061    1  0 Aug04 ?  00:00:00 /usr/sbin/atd -f
root 1063    1  0 Aug04 ?  00:00:00 /usr/bin/lxcs /var/lib/lxcs/
syslog 1069    1  0 Aug04 ?  00:00:00 /usr/sbin/syslogd -n
root 1074    1  0 Aug04 ?  00:00:00 /usr/sbin/sshd -D
root 25393  1074  0 01:31 ?  00:00:00 \_ sshd: ubuntu [priv]
ubuntu 25453  25393  0 01:31 ?  00:00:00 \_ sshd: ubuntu@pts/0
ubuntu 25454  25453  0 01:31 pts/0 00:00:00 \_ bash
ubuntu 25509  25454  0 01:35 pts/0 00:00:00 \_ ps --forest -eaf
root 1081    1  0 Aug04 ?  00:00:01 /usr/lib/snapd/snapd
root 1118    1  0 Aug04 ?  00:00:00 /sbin/mdadm --monitor --pid-file /run/mdadm/monitor.pid --daemon1
[...]

```

- Process Creation

- o After creating a child, the parent continues executing
 - Child process cannot run without its parent
- o But at any point, even right away, it can wait for the child's completion
- o Child can be:
 - A complete clone of the parent, or rather an exact copy of the address space
 - Be an entirely new program
- o POSIX Standard:
 - UNIX (Linux)
 - Darwin (MacOS, iOS, tvOS, watchOS)

- Fork() system call

- A system call that creates a new process
 - A really thin layer over the more complicated “clone” system call
 - “fork”, however, is kept as a system call for backward compatibility reasons
 - “man 2 fork”
- Child is almost an exact copy of parent, except for:
 - PID (two process cannot have the same ID)
 - PPID (parent cannot be its grandparent)
 - Resource utilization (set to 0 as it has just started)
- After call, parent continues executing and the child begins executing
- **The confusing part: “fork()” returns an integer:**
 - **Fork() returns 0 to the child**
 - **Fork() returns the child’s PID to parent**
 - In the case of an error (process table is full), fork() returns -1

fork(): The basic example

The basic use of fork()

```
returnValue = fork();
if (returnValue < 0) {
    // Manage the error
    print("Error: Can't fork!");
} else if (returnValue == 0) {
    // Child code
    print("I am the child and my pid is %d", getpid());
    while (1==1); // I just don't want it to terminate
} else {
    // Parent code
    print("I am the parent and the pid of my child is %d",
        returnValue);
    while (1==1); // I just don't want it to terminate either
}
```

- Simplified version of [fork_example1.c](#)
- Note: Errors cases should always be handled (but perhaps not for printf?)

fork(): Second example

What does the [following code](#) print?

Second fork() example

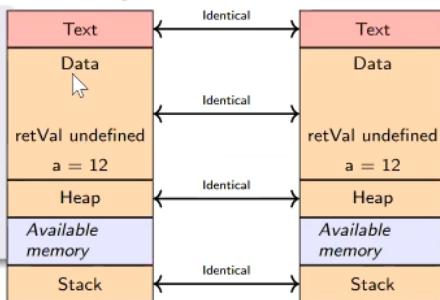
```
int a = 12;
 retVal = fork();
if (retVal) {
    // The PARENT (or error)
    sleep(5); // Ask the OS to put me in the WAITING state for 5s
    printf("a = %d", a); // Display the value of a
    while (1); // Loop forever
} else {
    // The CHILD
    a += 3;
    while (1); // Loop forever
}
```

- In this case, you would see the value 12 or 15 on the screen

fork(): Second example

Right after `fork()` and **before** the assignment to `retVal`

```
int a = 12;
retVal = fork();
if (retVal) {
    sleep(5);
    printf("a = %d", a);
    while (1);
} else {
    a += 3;
    while (1);
}
```



fork(): Second example

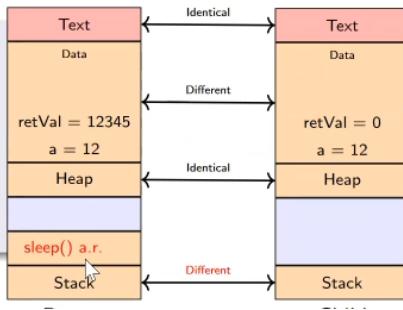
There are now two processes executing independently (assuming no error)

In the Parent

First an activation record is created because of

the sleep call.

```
int a = 12;
 retVal = fork();
if (retVal) {
    sleep(5);
    printf("a = %d", a);
    while (1);
} else {
    a += 3;
    while (1);
}
```



- Both processes coexist independently
 - o The code is executed **independently** in the parent and child
 - o The segment for the parent has **nothing to do** with the child
 - o The stack of the parent has **nothing to do** with the stack of the child
 - o The heap of the parent has **nothing to do** with the heap of the child

fork(): Second example tweaked

No infinite loop, print after the if statement

```
int a = 12;
retVal = fork();
if (retVal) {
    sleep(5);
} else {
    a += 3;
}
printf("%d", a);
```

What does this program print? Any ideas?...

- The program prints 1512

fork() can be confusing

How many times does this program print "Hello" ([source code](#))?

```
fork();  
printf("Hello");  
fork();  
printf("Hello");
```

Show of hands?

- This will print 6 times
-

fork(): “crazy example” (just for practice)

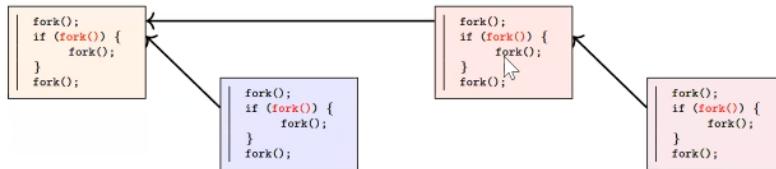
Question: How many processes does this C program create?

```
fork();
if (fork()) {
    fork();
}
fork();
```

- The above code is (likely) not useful
 - Note the typical C coding style: call `fork()` and execute the `if` clause if the return value is non-zero
- This is just a complex example that you can use to double-check that you truly understand `fork()`
- Let's go through the solution...

fork(): “crazy example”

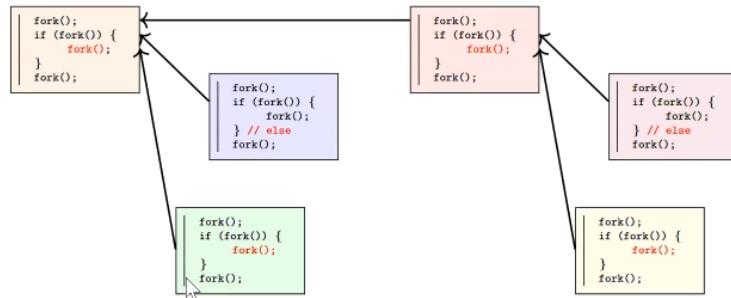
The second `fork()` (the one in the `if` condition) is executed



We now have **4** processes. Two of them return from the second `fork()` with non-zero, and two of them with zero

fork(): “crazy example”

The third fork() gets executed by only two processes, thus creating only two new processes:



We now have **6** processes. ALL of them call the last fork(), meaning that we have **12 processes** in total

- Now we can better understand the fork bomb that was run on the VM a week back
- Int main() {
 - o While (1) {
 - Fork()
 - o };
- }
- Best way to prevent forkbombs from killing system is to limit the number of processes a user can use via “ulimit -u <count>”
 - o Calls to “fork()” will fail before process table is full
- The Exec*() family:
 - o “man 3 exec:
 - o Exec, execp, execle, execv, execvp, execvpe – execute a file
 - o These are all variations of the execve system call and typically one says “exec”
 - o The system call **replaces the process image** (such as the process address space) **by that of a specific program** (stored on disk as an executable)
 - o Essentially, you give “exec”:
 - A path to an executable
 - Arguments to be passed to that executable
 - Possibly a set of environment variables
 - o The call to exec() never returns
 - Unless there is an error such as the executable cannot be executed

The exec*() Family: Example 1



```
int main(int argc, char *argv[]) {
    char* const arguments[] = {"ls", "-l", "/tmp", NULL};
    execv("/bin/ls", arguments);
    printf("This never gets executed...\n");
    return 0;
}
```

The exec*() Family: Example 4

To spawn a new process: combine with fork()

```
if (fork() == 0) {
    // Child
    char* const arguments[] = {"ls", "-l", "/tmp", NULL};
    execv("/bin/ls", arguments);
} else {
    // Parent
    while (1);
}
```

The exec*() Family: Examples 2 and 3

```
char* const arguments[] = {"ls", "-l", "/tmp", NULL};  
if (execv("Let's try something fun", arguments) == -1) {  
    fprintf(stderr, "Error: %s\n", strerror(errno));  
}
```

[exec_example2.c](#)

```
char* const arguments[] = {"ls", "-l", "/tmp", NULL};  
if (execv("/tmp", arguments) == -1) {  
    fprintf(stderr, "Error: %s\n", strerror(errno));  
}
```

[exec_example3.c](#)

They're dead... but alive!

Let's have a look at the processes list while the last example is running...

```
% ps f  
  PID  TTY      STAT      TIME      COMMAND  
 9129  pts/18    Ss        0:00      bash  
10418  pts/18    R+        0:03      \_ ./exec_example4  
10419  pts/18    Z+        0:00      \_ [ls] <defunct>
```

Defunct: from Latin *defunctus*, i.e., dead

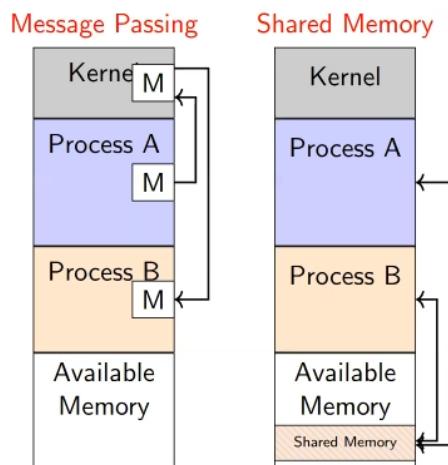


- When a child process terminates, it remains as a **zombie** in the **terminated** state
 - o Recall that in the Process Lifecycle Diagram, we had a terminated state, which some of us might have thought was useless
- Rationale: the parent process may want to know about the status of a child process:
 - o i.e., Run this first program and when it is finished, run that one...
 - o therefore, the parent may want to wait for its child's completion
- The OS keeps zombies around for this purpose

- Zombies do not use hardware resources
 - But each zombie uses a slot in the process table!
 - The Process Table may thus eventually fill up due to Zombies and thus cause “fork()” to fail
- A zombie lingers on until:
 - **Its parent has acknowledged termination**
 - **Its parent dies**
- A process terminates itself with the “exit()” system call, which takes as argument an integer called the process **exit/return/error value/code**
- All resources of the process are then deallocated by the OS (i.e., memory, open files, I/O buffers, etc.)
- Note that the termination of a process can terminate another process
- This is done using **signals** and the “kill()” system call
- Communicating Processes?
 - So far we have seen **independent** processes
 - Each process runs code independently
 - Parents are aware of their children, and children are aware of their parents. They do not interact, however aside from the ability to wait for a process to terminate
 - We often need processes to **cooperate** however:
 - To share information such as access to common data
 - To speed up computation such as using multiple cores
 - For the sake of convenience such as some applications are naturally implemented as sets of interacting process

Two Broad Communication Models

- Example: Process A needs to communicate with Process B



Message Passing	Shared Memory
😢 Requires data copies (user space → kernel space → user space)	😊 Fast because no memory copy is needed
😢 High overhead because one system call per operation (send and receive)	😊 Low overhead because system calls only to set up the shared memory region (after that it's just memory reads and writes)
😢 Cumbersome for developers: code will be sprinkled with send and receive everywhere	😊 Easy for developers: Just read/write to shared memory region
😊 Preserves memory protection / isolation across processes	😢 Violates the principle of memory protection between processes, which is dangerous / bug-prone (more on this in Synchronization module)

- All OSes provide system calls for Shared Memory
- All OSes provide system calls for Message passing
- There are many variations in APIs and implementations